

Introduction

Machine Learning

Definition (Machine Learning)

Spicy statistics/signal processing. Methods, algorithms, theorems, etc. to make inferences from data.

Three (main) types of learning:

- **Supervised learning.** Given data set of points and labels, find a prediction function that, given any point, finds a corresponding label.
 - Useful in: a lot. Analytics, autopilot, security, face recognition, etc.
- **Unsupervised learning.** Given data set of *only* points (*labels hidden to us*), find a prediction function.
 - Useful in: cases where you only need similarity between datapoints, or have too much training data to label by hand. Collaborative filtering, data visualization, etc.
- **Reinforcement learning.** Given *no* data, but rather an environment, collect data from environment and learn to do a task optimally. Will skip this in this course.

So What Are We Actually Studying?

Mainly results/intuition/algorithms. Very few proofs, if any. (Take an actual class!)

Course organization: Data-Driven. Two types of data:

- **Unstructured Data.** (Useful for e.g. analytics, security, other traditional statistical tasks.)
 - Probabilistic classifiers.
 - Linear/logistic regression.
 - Regularization.
 - Dimensionality reduction.
 - Kernels.
 - Clustering.
 - Decision trees.
 - Boosting and ensembles.
- **Structured Data.** (Useful for e.g. image/text/audio processing.)
 - Sparse vector/matrix recovery.
 - Sparse vector/matrix completion.
 - Sparse/robust dimensionality reduction.
 - Dictionary learning.
- **But wait, where are neural networks?**
 - Last few lectures: neural networks for structured/unstructured data.

Different Spaces We Talk About

- \mathcal{X} is **input space**.
 - For a list of numbers of length d , $\mathcal{X} = \mathbb{R}^d$.
 - For a 128×128 grayscale image, $\mathcal{X} = \mathbb{R}^{128 \times 128}$.
- \mathcal{Y} is **output space**.
 - For a classification task into k classes, $\mathcal{Y} = \{1, 2, \dots, k\} = [k]$.
 - For a regression task (predicting a continuous label), $\mathcal{Y} = \mathbb{R}$.
 - For an unsupervised representation learning task (embedding data into \mathbb{R}^d), $\mathcal{Y} = \mathbb{R}^d$.
- \mathcal{Z} is **data space**.
 - In a supervised problem, $\mathcal{Z} = \mathcal{X} \times \mathcal{Y}$.
 - In an unsupervised problem, $\mathcal{Z} = \mathcal{X}$.
- \mathcal{H} is a **hypothesis class** ; a class of functions $\mathcal{X} \rightarrow \mathcal{Y}$ we should care about

Technical note:

In general, we assume that these spaces have lots of structure: a rich topology, a measure on the Borel σ -algebra, and so on. This is fine because we can usually embed these spaces in \mathbb{R}^d for large enough d and use the structure of \mathbb{R}^d to give a structure to \mathcal{X} and \mathcal{Y} . The question of \mathcal{H} is a bit harder and requires functional analysis.

Notation/Vocabulary Dump

Definition, Assumptions (Data)

- The data point x_i is a random variable which has distribution μ_X on \mathcal{X} . Has density p_X .
 - Use x to denote arbitrary data point with distribution μ_X .
- The label y_i is a random variable which has distribution μ_Y on \mathcal{Y} . Has density p_Y .
 - Use y to denote arbitrary label with marginal μ_Y .
- The data z_i is a random variable which has distribution μ on \mathcal{Z} . Has density p_Z .
 - In a supervised learning problem, $z_i = (x_i, y_i)$.
 - In an unsupervised learning problem, $z_i = x_i$ and $\mu = \mu_X$.
 - Use z to denote arbitrary data with distribution μ .
- The data $\{z_i\}_{i=1}^n$ are independent and identically distributed (i.i.d.) as μ .

Definition (Sample)

A **sample** is n data points: $s = \{z_1, \dots, z_n\} \sim \mu^n$. Has density p_S .

Sometimes we use the notation $s_X = \{x_1, \dots, x_n\}$ or $s_Y = \{y_1, \dots, y_n\}$.

Notation/Vocabulary Dump

Definition (Loss)

A **loss function** $\ell: \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$ quantifies how bad our prediction is. High loss = bad.

Example (Classification Loss)

A popular classification loss is **multi-class cross-entropy loss**:

$$\ell(y_{\text{true}}, y_{\text{pred}}) = - \sum_{i=1}^k y_{\text{true},i} \log(y_{\text{pred},i})$$

Example (Squared-Error Loss)

A popular regression loss is **squared error loss**:

$$\ell(y_{\text{true}}, y_{\text{pred}}) = \|y_{\text{true}} - y_{\text{pred}}\|_2^2$$

Definition (Loss of A Hypothesis)

Given a hypothesis h , we can define the loss $\ell_h: \mathcal{Z} \rightarrow \mathbb{R}$.

In supervised learning, $\ell_h(z) = \ell_h((x, y)) = \ell(y, h(x))$.

Also define the loss over a sample s using L_s or distribution ν using L_ν :

$$L_s(h) = \mathbb{E}_{z \sim \text{Uni}(s)} [\ell_h(z)] = \frac{1}{n} \sum_{i=1}^n \ell_h(z_i)$$

$$L_\nu(h) = \mathbb{E}_{z \sim \nu} [\ell_h(z)]$$

Machine Learning Problem

A machine learning *problem* is specified by the following:

- An **input space** \mathcal{X} .
- An **output space** \mathcal{Y} .
- An **loss function** ℓ .

(The hypothesis class \mathcal{H} is a function of the algorithm. If we have prior knowledge what \mathcal{H} *should* be, we can choose an algorithm that gets this \mathcal{H} .)

How do we come up with solutions? Here's a standard workflow.

- ➊ **Collect a sample** s . Split it into a training set s_{train} and a validation set s_{val} .
- ➋ **Run a machine learning algorithm** $h = A(s_{\text{train}})$.
- ➌ **Evaluate the output on the validation set.** Compute $L_{s_{\text{val}}}(h)$.
- ➍ **Is your accuracy $L_{s_{\text{val}}}(h)$ good enough?**
 - If yes, you're done! Store the model somewhere for use, or keep improving it.
 - If not, either collect more data, improve data quality, or change the algorithm.

Data collection/quality is out of scope, but we can tune algorithms.

- Tune algorithm parameters (these are **hyperparameters** and usually set by hand). See k -fold cross-validation for how to efficiently put this into the workflow.
- Change the algorithm completely.